

From Accurate Forecasts to Feasible Plans: Feasibility-Constrained Model Selection for Operational Demand Planning

Author Name Placeholder

Affiliation Placeholder

City, Country Placeholder

email.placeholder@example.com

Abstract

Operational demand planning increasingly relies on predictive models, but operations execute planning signals rather than forecast accuracy metrics. A numerically accurate forecast can still produce volatile inventory targets, frequent switching, or execution violations when downstream systems cannot absorb abrupt plan changes. This paper formulates forecast-driven demand planning as feasibility-constrained model selection. Candidate forecasts are converted into executable planning signals and evaluated using WAPE, inventory exposure, planning volatility, switching burden, execution violations, normalized planning loss, greedy-to-DP gap, and oracle gaps. DataCo execution records provide data-informed execution-risk scenarios, Favorita retail demand data provide the main forecast-to-plan proof of concept, and M5 retail data test hierarchy and intermittent-demand robustness. Across 123 paper-facing evaluation groups, the accuracy-first deployable strategy differs from the operational-loss-optimal deployable strategy in 91.1% of cases. The accuracy-first planning-execution gap rate is 15.4%, compared with 2.7% for the best deployable operational-loss strategy. Favorita results show ranking shifts under severe execution-risk scenarios, while M5 results show that planning grain and intermittency change the preferred deployable mechanism. The contribution is not a new forecasting model. It is an interpretable operations decision layer showing that forecast accuracy should be treated as a diagnostic metric, while deployment should be evaluated by feasibility-adjusted planning utility.

Keywords

Demand planning, Forecast-to-plan conversion, Operational feasibility, Planning stability, Execution risk.

1 Introduction

Industrial demand-planning systems increasingly rely on predictive models to support replenishment, inventory positioning, labor planning, transportation commitments, and capacity allocation. Recent reviews show that artificial intelligence and machine learning are becoming more visible in supply-chain planning and demand-planning practice (Walter et al. 2025; Culot et al. 2024; Vlachos and Reddy 2025). Forecasting benchmarks and time-series foundation models further show that predictive capability is advancing rapidly (Makridakis et al. 2022; Das et al. 2024; Ansari et al. 2024). These advances create value, but they also sharpen an operations-management problem: better forecast accuracy does not automatically mean better deployment quality.

Operations execute planning signals, not forecast accuracy metrics. A planner cannot execute WAPE or RMSE. The operation executes an inventory target, a replenishment order, a staffing plan, a production target, or a transportation commitment. When a forecast-driven plan changes sharply, downstream systems may need supplier coordination, warehouse rebalancing, labor rescheduling, or enterprise planning-system updates. These changes can be costly even when the forecast itself is numerically accurate.

This creates a planning-execution gap. Forecasting models may adapt daily or weekly, while procurement lead times, replenishment workflows, planner governance, and execution infrastructure may absorb only gradual change. A forecast-driven plan can therefore create planning volatility, model-switching burden, and execution-capacity violations. Standard forecasting evaluation often emphasizes WAPE, MAE, RMSE, or related accuracy metrics, while deployment requires feasibility-adjusted planning utility.

This paper studies feasibility-constrained model selection for operational demand planning. The central thesis is that forecasting-driven strategies should be selected by operational feasibility as well as prediction accuracy. Candidate

forecasts are converted into planning signals, evaluated against inventory exposure and execution constraints, and compared through deployable decision-layer policies. The empirical contribution is to measure how often and by how much accuracy-first selection differs from operational-loss-optimal deployable selection.

The study uses three public-data roles. DataCo supply-chain execution records are used to anchor execution-risk scenarios from observed late-delivery heterogeneity. Favorita daily retail demand data provide the main forecast-to-plan proof of concept. M5 retail data provide hierarchy and intermittent-demand robustness. These datasets are used to evaluate whether the planning-execution gap remains visible under different operational conditions; they are not used to claim that one forecasting model universally dominates.

The paper contributes to IEOM-oriented operations research in four ways. First, it defines a forecast-to-plan evaluation layer that reports forecast accuracy separately from operational planning utility. Second, it operationalizes the planning-execution gap through execution violation rate and planning-execution gap rate. Third, it compares accuracy-first, ensemble, smoothing, greedy, dynamic programming, and budgeted dynamic programming deployable strategies. Fourth, it validates the tradeoff using DataCo-informed scenarios, Favorita proof-of-concept results, and M5 hierarchy/intermittency robustness.

1.1 Objectives

The specific objectives of this study are:

1. Define a forecast-to-plan evaluation layer that measures operational feasibility in addition to WAPE.
2. Operationalize the planning-execution gap using execution violation rate and planning-execution gap rate.
3. Use DataCo execution records to construct data-informed execution-risk scenarios.
4. Compare accuracy-first, ensemble, smoothing, greedy, DP, and budgeted-DP deployable planning strategies.
5. Validate the planning-execution gap across Favorita and M5 and translate the findings into practical planning recommendations.

2 Literature Review

2.1 AI-enabled Demand Planning and Forecasting Benchmarks

Recent supply-chain research shows increasing interest in AI-enabled demand planning. Walter et al. (2025) review artificial intelligence in demand planning and identify AI tools, supply-chain applications, and digital supply-chain impacts as major knowledge clusters. Culot et al. (2024) review empirical AI studies in supply-chain management and emphasize the need to connect AI adoption to operational research directions. Vlachos and Reddy (2025) similarly review machine learning in supply-chain management and identify planning, execution, and control as areas where machine learning can support decision making.

Forecasting benchmarks and time-series foundation models show the same trend from a forecasting perspective. The M5 competition demonstrates the scale and hierarchy of modern retail forecasting benchmarks (Makridakis et al. 2022). Recent time-series foundation models, including decoder-only forecasting models and pretrained probabilistic time-series models, show that forecasting pipelines are moving toward broader zero-shot and cross-domain capability (Das et al. 2024; Ansari et al. 2024). These studies motivate the forecasting layer of the present paper, but they do not remove the need to evaluate whether forecast-driven planning signals are operationally executable.

2.2 Forecast-to-decision and Predict-then-optimize Research

Forecast-to-decision research is directly relevant because it explains why prediction accuracy and downstream decision quality can differ. Elmachoub and Grigas (2022) show that the standard predict-then-optimize workflow can

be improved by training predictions toward decision loss rather than prediction loss alone. Mandi et al. (2024) survey decision-focused learning and describe how machine learning and constrained optimization can be integrated to improve decision quality under uncertainty. Jia et al. (2025) propose a scenario predict-then-optimize method for data-driven online inventory routing, showing that downstream routing and replenishment decisions require more than point forecast accuracy.

The present paper is related to this stream, but its purpose is different. It does not introduce a new end-to-end learning algorithm. It proposes an interpretable forecast-to-plan decision layer for operational demand planning, where existing candidate forecasts are evaluated by inventory exposure, planning volatility, switching burden, and execution-capacity violations.

2.3 Inventory Implications and Planning Stability

Inventory theory provides the operational foundation for holding and shortage costs. Silver et al. (1998) and Axsater (2015) describe inventory planning as a cost and service-level problem rather than a pure prediction problem. Forecast accuracy metrics remain important, but Hyndman and Koehler (2006) show that forecast-error measures have known limitations and must be interpreted carefully across series and contexts. Forecast combination is also relevant because weighted or averaged forecasts can reduce dependence on a single model (Bates and Granger 1969).

Modern forecasting candidate sets often include nonlinear machine-learning methods such as random forests, XG-Boost, and LightGBM (Breiman 2001; Chen and Guestrin 2016; Ke et al. 2017). These models can improve forecast accuracy, but model deployment also creates governance, monitoring, and switching burden. From an operations perspective, Forrester (1961) shows that industrial systems can amplify instability when material and information flows respond with delays. Advanced planning systems similarly separate forecasting, planning, coordination, and execution layers (Stadtler et al. 2015). This separation motivates explicit measurement of planning volatility and execution-capacity violations.

Rolling-horizon decision theory also supports finite-horizon selection logic because a locally attractive action can be undesirable once future transitions are considered (Sethi and Sorger 1991). This perspective is consistent with the paper's use of greedy-to-DP comparisons to measure myopic planning suboptimality.

2.4 Research Gap

Existing studies improve forecasts, combine forecasts, or optimize downstream decisions. Fewer studies provide a practical decision layer that compares forecasting-driven strategies by inventory exposure, planning volatility, switching burden, and execution-capacity violations. This paper addresses that gap by defining planning-execution gap rate and feasibility-adjusted normalized planning loss. The goal is not to replace forecasting metrics, but to position WAPE as a diagnostic measure alongside operational feasibility.

3 Methods

The proposed method follows the IEOM sample-paper pattern of a visual workflow followed by stepwise algorithmic logic. Figure 1 gives the framework overview. Candidate forecasts are generated first, but the decision layer evaluates their implied planning signals before choosing a deployable planning strategy. This framing follows the forecast-to-decision logic that downstream objectives may differ from forecast-error objectives (Elmachtoub and Grigas 2022; Mandi et al. 2024).

3.1 Feasibility-constrained Planning Model

Let i index a planning unit, t index a planning period, and m index a candidate forecasting model. The candidate forecast is $\hat{y}_{i,t}^m$, and the executable planning signal is $p_{i,t}^m$. The planning signal is created through a forecast-to-plan




Figure 1: Feasibility-constrained planning framework.

conversion rule:

$$p_{i,t}^m = g(\hat{y}_{i,t}^m, \text{policy}_t). \quad (1)$$

In the current implementation, $p_{i,t}^m$ is an inventory target implied by the forecast and safety-stock policy.

Inventory exposure is measured through holding and shortage terms:

$$IC_{i,t}^m = h[p_{i,t}^m - y_{i,t}]^+ + b[y_{i,t} - p_{i,t}^m]^+, \quad (2)$$

where h is holding cost, b is shortage cost, $y_{i,t}$ is realized demand, and $[x]^+ = \max(x, 0)$.

Planning volatility is the absolute movement in the planning signal:

$$PV_{i,t}^m = |p_{i,t}^m - p_{i,t-1}^m|. \quad (3)$$

Execution capacity $C_{i,t}$ is the amount of plan movement the operation can absorb in period t . The execution violation is:

$$EV_{i,t}^m = [|p_{i,t}^m - p_{i,t-1}^m| - C_{i,t}]^+. \quad (4)$$

The paper-facing planning-execution gap rate for strategy s is:

$$PEG(s) = \frac{1}{T} \sum_{t=1}^T \mathbb{I}\{EV_t(s) > 0\}. \quad (5)$$

The normalized planning loss is:

$$L(s) = \beta_{\text{inventory}} \cdot \text{NIC}(s) + \lambda_{\text{volatility}} \cdot \text{NVC}(s) + \lambda_{\text{execution}} \cdot \text{NEC}(s) + \lambda_{\text{switch}} \cdot \text{NSC}(s), \quad (6)$$

where NIC, NVC, NEC, and NSC denote normalized inventory, volatility, execution, and switching components. WAPE is reported as diagnostic forecast accuracy. The main operational objective excludes direct forecast-error terms because inventory exposure and feasibility terms capture the deployable consequences of forecast-driven planning decisions.

3.2 Feasibility-constrained Selection Procedure

The selection procedure has ten steps.

Step 1. Define planning units and candidate forecasting models. The planning units are store-family or item-store style series, and candidate models include transparent baselines, machine-learning candidates where available, ensembles, smoothing policies, and decision-layer selectors.

Step 2. Generate candidate forecasts. Each candidate model produces a forecast for each planning unit and period

using only information available before the forecast timestamp.

Step 3. Convert each forecast into a planning signal. The forecast is translated into an inventory planning signal using the policy function $g(\cdot)$.

Step 4. Compute inventory exposure. Holding and shortage terms are computed from the candidate planning signal and realized demand for evaluation.

Step 5. Compute planning volatility. Each candidate signal is compared with the previous executed plan to measure plan movement.

Step 6. Compute execution violation. The plan movement is compared with execution capacity $C_{i,t}$ to determine whether the period creates a planning-execution gap.

Step 7. Compute switching burden. For hard model selection, a switching indicator is recorded when the selected model differs from the previous period's model.

Step 8. Compute normalized operational loss. Inventory, volatility, execution, and switching components are normalized within the relevant evaluation group and combined using the stated planning weights.

Step 9. Select the deployable planning action. The decision layer selects a candidate plan using an accuracy-first, ensemble, smoothing, greedy, DP, or budgeted-DP policy.

Step 10. Validate against oracle benchmarks. Deployable strategies are compared with non-deployable oracle diagnostics to separate selection suboptimality from forecast uncertainty.

3.3 Greedy and Dynamic Programming Selectors

The greedy selector minimizes current-period operational cost only. Let $c_t(m_{t-1}, m)$ denote the validation-derived stage cost of switching from previous model m_{t-1} to candidate model m at time t . The greedy selector chooses:

$$m_t^G = \arg \min_{m \in \mathcal{M}} c_t(m_{t-1}, m). \quad (7)$$

This policy is interpretable but myopic because it ignores future switching and execution consequences.

The finite-horizon DP selector accounts for future operational costs:

$$V_t(m_{t-1}) = \min_{m \in \mathcal{M}} \{c_t(m_{t-1}, m) + V_{t+1}(m)\}, \quad (8)$$

with terminal condition:

$$V_{T+1}(\cdot) = 0. \quad (9)$$

The greedy-to-DP gap measures myopic selection suboptimality. DP is included to quantify the value of planning ahead over one-step operational selection.

3.4 Budgeted DP Selector

Some planning organizations impose governance limits on model changes. The budgeted DP selector adds:

$$\sum_{t=2}^T \mathbb{I}\{m_t \neq m_{t-1}\} \leq K. \quad (10)$$

The DP state includes the previous model and the number of switches already used. Transitions that exceed the switch budget are discarded. If missing candidate availability leaves no feasible path, the implementation uses a flagged incumbent-stays fallback: it chooses the first period's lowest-cost model and holds that incumbent when available. This fallback is budget-compliant when the incumbent model is available throughout the horizon and is marked diagnostically if availability fails.

3.5 Ensemble and Smoothing Policies

The Simple Ensemble averages candidate forecasts without using operational-loss information. The Operational-Loss Ensemble uses validation normalized planning loss to assign nonnegative weights. These policies test whether soft forecast composition can reduce abrupt execution burden while retaining useful accuracy.

The smoothing policy changes the executed planning signal, not the underlying forecast:

$$p_t = \alpha p_t^{\text{candidate}} + (1 - \alpha) p_{t-1}. \quad (11)$$

Lower α values reduce execution burden by adapting gradually, but they can increase inventory exposure if the plan responds too slowly to demand changes.

3.6 Oracle Benchmarks and Validation Gaps

The Realized-Inventory Oracle DP is non-deployable. It uses realized test-period inventory outcomes over the same forecast candidates and is used only to quantify the value of unavailable inventory-outcome information. It is not a perfect-forecast oracle.

The Perfect-Demand Oracle uses realized demand as the forecast signal. It is also non-deployable and can still create volatility or execution penalties because perfect demand tracking may move faster than execution capacity. The metric `gap_to_dp_oracle` measures selection suboptimality relative to the realized-inventory DP benchmark. The metric `gap_to_perfect_oracle` measures the total gap to a realized-demand reference, including both forecast uncertainty and feasibility effects.

Table 1 links each method component to its evidence and operational implication.

Table 1: Method components, validation evidence, and operational implication.

Method component	What it tests	Paper-facing evidence	Operational implication
Forecast-to-plan conversion	Whether forecasts become materially different planning signals	Accuracy-first mismatch rate and Favorita frontier	Evaluate plans, not forecasts alone.
Planning-execution gap rate	Whether plans exceed execution capacity	PEG rate: 15.4% accuracy-first vs. 2.7% best deployable	Monitor feasibility as a deployment metric.
DataCo scenarios	Whether empirical execution-risk regimes reorder strategies	Baseline-to-severe Favorita ranking shifts	Use scenario tests before rollout.
Greedy selector	Whether one-step selection is sufficient	Greedy-to-DP gap	Greedy policies can be myopic.
Finite-horizon DP	Value of planning ahead under transition costs	Greedy-to-DP gap and DP oracle gap	Account for future switching and execution burden.
Budgeted DP	Governance limit on model switching	Hard switch budget and flagged fallback semantics	Use when model changes require approval.
Smoothing policy	Gradual operational adaptation	Lower execution burden with inventory tradeoff	Adapt slowly when infrastructure is constrained.
Ensemble policy	Soft composition across models	Lower execution burden in Favorita severe scenario	Reduce model-specific plan jumps.
Oracle benchmarks	Value of unavailable future information	<code>gap_to_dp_oracle</code> and <code>gap_to_perfect_oracle</code>	Separate selection suboptimality from forecast uncertainty.
M5 robustness	Sensitivity to planning grain and intermittency	Best method changes across M5 slices	Avoid one-size-fits-all deployment rules.

4 Data Collection and Experimental Design

The empirical design uses DataCo, Favorita, and M5 for distinct operational roles. Table 2 summarizes these roles. DataCo is not treated as an inventory-control dataset; it is used to transfer empirical heterogeneity in execution risk into scenario design. Favorita is the main forecast-to-plan proof-of-concept dataset. M5 is the hierarchy and intermittent-demand robustness dataset.

DataCo execution-risk anchor. DataCo contains supply-chain fulfillment records with late-delivery outcomes. The analysis extracts late-delivery anchors from context-level percentiles. The current generated anchors range from 0.535

Table 2: Dataset roles and experimental purpose.

Dataset	Role	Data used	Purpose
DataCo	Execution-risk anchor	Late-delivery outcomes by operational context	Defines data-informed execution-risk scenarios; not treated as an inventory-control dataset.
Favorita	Main proof of concept	Daily store-family retail demand panel	Tests forecast-to-plan conversion and deployable strategy ranking under execution-risk scenarios.
M5	Robustness analysis	Hierarchical and intermittent retail demand	Tests whether the planning-execution gap persists across planning grain and demand sparsity.

at the low-risk percentile to 0.901 at the severe-risk percentile. Under the additive clipped scenario mapping, these anchors produce execution weights from 0.318 to 0.500. These are sensitivity weights, not exact dollar costs for Favorita or M5.

Favorita proof of concept. Favorita is a daily retail demand-planning panel from Corporacion Favorita. The pipeline treats each store-family pair as a planning unit and uses promotion, holiday, oil-price, transaction, and store-context information where available at planning time. Candidate forecasts include transparent baselines and machine-learning candidates such as LightGBM and XGBoost when dependencies are available (Ke et al. 2017; Chen and Guestrin 2016). The purpose is to test whether forecast-to-plan conversion changes deployable strategy ranking.

M5 robustness design. M5 contains 3,049 items across 10 stores, and the pipeline samples item-store series for quick and medium robustness runs (Makridakis et al. 2022). The IEOM submission uses M5 to evaluate item-store, department-store, and category-store planning grains and to test intermittent-demand buckets. This design checks whether the planning-execution gap persists when demand is sparse or aggregated.

Walmart and Rossmann are not completed experiments in this IEOM draft. They can be used later as optional future robustness panels after the DataCo-Favorita-M5 story is finalized.

5 Results and Discussion

5.1 Numerical Results

Quantifying the planning-execution gap. The strongest evidence for the planning-execution gap is the mismatch between the WAPE-best deployable strategy and the operational-loss-best deployable strategy. Across 123 paper-facing evaluation groups, the accuracy-first deployable method differs from the operational-loss-optimal deployable method in 91.1% of cases. The accuracy-first planning-execution gap rate is 15.4%, compared with 2.7% for the best deployable operational-loss strategy. Table 3 summarizes the core numerical evidence.

Ranking shifts under DataCo-informed execution risk. Favorita results show that Global Best remains a strong accuracy-first reference. In the baseline scenario, Global Best has WAPE 0.109 and normalized total loss 1.550, while Operational Ensemble has WAPE 0.119 and normalized total loss 1.422. Under the severe DataCo-informed execution-risk scenario, Global Best remains accurate but has normalized total loss 2.000. Simple Ensemble and Operational Ensemble have lower normalized losses of 1.473 and 1.496, respectively. This does not imply a universal winner. It shows that deployable ranking depends on execution-risk regime.

Robustness under hierarchy and intermittency. M5 does not simply replicate Favorita. Table 4 shows that Operational Ensemble is strongest at department-store and category-store grains, while Smoothed Alpha 0.25 is strongest at the item-store grain and across low-, medium-, and high-intermittency buckets. This strengthens the paper because the framework is not designed to make one method always win. Instead, it shows that the preferred deployable mechanism changes with operational context.

Table 3: Core numerical evidence for the planning-execution gap.

Metric	Current value	Interpretation
Accuracy-first mismatch rate	91.1% of 123 groups	WAPE-best deployable strategy often differs from operational-loss-best deployable strategy.
Mean accuracy-first loss gap	11.169	Average normalized-loss disadvantage of accuracy-first selection when it is not operationally best.
Accuracy-first PEG rate	15.4%	Fraction of periods with positive execution violation for accuracy-first strategies.
Best deployable PEG rate	2.7%	Same metric for the best deployable operational-loss strategy.
Greedy-to-DP gap	0.044	Measures myopic selection cost where greedy and DP are both available.
Mean gap to DP oracle	0.198	Selection suboptimality relative to non-deployable realized-inventory DP information.
Mean gap to perfect oracle	-8.936	Diagnostic total gap; negative value reflects that realized-demand tracking can create high execution burden.
DataCo execution-risk anchors	0.535 to 0.901	Late-delivery anchors used for low to severe execution-risk scenarios.

Table 4: M5 hierarchy and intermittency robustness summary.

Robustness slice	Best deployable method	WAPE	PEG rate	Norm. loss
Hierarchy: category-store	Operational Ensemble	0.159	0.063	1.261
Hierarchy: department-store	Operational Ensemble	0.162	0.100	1.177
Hierarchy: item-store	Smoothed Alpha 0.25	0.368	0.010	1.223
Intermittency: high	Smoothed Alpha 0.25	0.616	0.030	1.284
Intermittency: low	Smoothed Alpha 0.25	0.254	0.000	1.156
Intermittency: medium	Smoothed Alpha 0.25	0.367	0.001	1.168

Myopic selection and oracle gaps. The generated thesis summary reports a mean greedy-to-DP gap of 0.044, indicating that finite-horizon planning improves over one-step greedy selection where both are available. The mean gap to the realized-inventory DP oracle is 0.198, which measures selection suboptimality relative to non-deployable inventory-outcome information. The mean gap to the perfect-demand oracle is negative in the current combined outputs because a realized-demand signal can create high volatility and execution penalties; this confirms that the perfect-demand oracle is a diagnostic reference, not a deployable lower bound.

5.2 Graphical Results

Figure 2 plots the Favorita accuracy versus execution-penalty frontier. The figure shows that WAPE and execution penalty do not move together. Global Best is near the best WAPE region, but other deployable strategies produce much lower execution penalties. The operational implication is that the deployable frontier is not the WAPE-best point.




Figure 2: Favorita accuracy and execution penalty frontier.

Figure 3 shows normalized planning loss across DataCo-informed execution scenarios. As execution-risk weight changes, strategy ranking changes. This supports scenario testing as a deployment practice: the model selected for a flexible execution environment may not be the model selected for a constrained execution environment.

5.3 Proposed Improvements

Monitor planning-execution gap rate. The accuracy-first planning-execution gap rate is 15.4%, compared with 2.7% for the best deployable operational-loss strategy. This supports adding PEG rate to routine model-deployment monitoring.

Avoid WAPE-only model selection. The accuracy-first deployable method differs from the operational-loss-optimal deployable method in 91.1% of paper-facing evaluation groups. Planners should treat WAPE as diagnostic, not as the only deployment criterion.

Use execution-risk scenario testing. Favorita rankings change between baseline and severe DataCo-informed scenarios. Before rollout, planning teams should test candidate strategies under plausible execution-capacity regimes.

Use smoothing and ensembles selectively. Smoothing can reduce execution burden, but it can also increase inventory exposure when adaptation is too slow. Ensembles can reduce abrupt model-specific movements, but they may sacrifice WAPE. These methods should be selected by stated operational tradeoffs rather than by a universal rule.




Figure 3: Favorita normalized planning loss by DataCo-informed execution scenario.

Use DP or budgeted DP when switching governance matters. The greedy-to-DP gap shows that one-step selection can be myopic. Budgeted DP is useful when planners impose hard limits on model switching, approvals, or monitoring changes.

5.4 Validation

Baseline validation. Accuracy-first methods remain necessary baselines. The paper does not claim that feasibility-aware strategies always dominate. Instead, it shows when the accuracy-first strategy is not the operational-loss-optimal deployable strategy.

Oracle validation. The Realized-Inventory Oracle DP and Perfect-Demand Oracle are non-deployable diagnostics. The realized-inventory DP oracle measures model-selection suboptimality over the same forecast candidates. The perfect-demand oracle includes forecast uncertainty and can still create execution burden.

Scenario validation. DataCo-informed scenarios validate whether execution-risk weights alter deployable ranking. The baseline-to-severe Favorita shift shows that execution risk can reorder strategy preference.

Robustness validation. M5 validates the framework under hierarchy and intermittent demand. The best deployable mechanism changes across planning grain and intermittency bucket, which supports the paper's core tradeoff thesis rather than a single-method dominance claim.

6 Conclusion

This paper reframes operational demand planning as feasibility-constrained model selection. The planning-execution gap was quantified by comparing accuracy-first model selection with operational-loss-optimal deployable planning selection. Across 123 paper-facing evaluation groups, the accuracy-first deployable method differs from the operational-loss-best deployable method in 91.1% of cases. The accuracy-first planning-execution gap rate is 15.4%, while the best deployable operational-loss strategy has a gap rate of 2.7%.

Objective 1 was addressed by defining a forecast-to-plan evaluation layer that reports WAPE separately from inventory exposure, volatility, switching burden, and execution violations. Objective 2 was addressed by defining planning-execution gap rate. Objective 3 was addressed by using DataCo execution records to construct data-informed execution-risk scenarios. Objective 4 was addressed by comparing accuracy-first, ensemble, smoothing, greedy, DP,

and budgeted-DP deployable strategies. Objective 5 was addressed through Favorita and M5 validation and through evidence-linked planning recommendations.

The results show that execution-risk scenarios can change deployable strategy ranking and that hierarchy and intermit-tency can change the preferred mechanism. The practical implication is direct: planners should evaluate deployment feasibility before model rollout. Forecast accuracy remains important, but it should be interpreted as a diagnostic metric rather than as the sole deployment objective.

Acknowledgements

Acknowledgement placeholder. Replace this paragraph with funding, advisor, or institutional acknowledgements before submission if applicable.

Biography

Author biography placeholder. Replace this paragraph with a biography for each author before submission. IEOM guidance limits each biography to 250 words.

References

- Ansari, A. F., Stella, L., Turkmen, C., Zhang, X., Mercado, P., Shen, H., Shchur, O., Rangapuram, S. S., Pineda Arango, S., Kapoor, S., Zschiegner, J., Maddix, D. C., Wang, H., Mahoney, M. W., Torkkola, K., Wilson, A. G., Bohlke-Schneider, M. and Wang, Y., Chronos: Learning the language of time series, *Transactions on Machine Learning Research*, October 2024.
- Axsater, S., *Inventory Control*, 3rd Edition, Springer, 2015.
- Bates, J. M. and Granger, C. W. J., The combination of forecasts, *Operational Research Quarterly*, vol. 20, no. 4, pp. 451-468, 1969.
- Breiman, L., Random forests, *Machine Learning*, vol. 45, no. 1, pp. 5-32, 2001.
- Chen, T. and Guestrin, C., XGBoost: A scalable tree boosting system, *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 785-794, 2016.
- Culot, G., Podrecca, M. and Nassimbeni, G., Artificial intelligence in supply chain management: A systematic literature review of empirical studies and research directions, *Computers in Industry*, vol. 162, article 104132, 2024.
- Das, A., Kong, W., Sen, R. and Zhou, Y., A decoder-only foundation model for time-series forecasting, *Proceedings of the 41st International Conference on Machine Learning*, PMLR, vol. 235, pp. 10148-10167, 2024.
- Elmachtoub, A. N. and Grigas, P., Smart predict, then optimize, *Management Science*, vol. 68, no. 1, pp. 9-26, 2022.
- Forrester, J. W., *Industrial Dynamics*, MIT Press, Cambridge, MA, 1961.
- Hyndman, R. J. and Koehler, A. B., Another look at measures of forecast accuracy, *International Journal of Forecasting*, vol. 22, no. 4, pp. 679-688, 2006.
- Jia, M., Schrottenboer, A. H. and Chen, F., Scenario predict-then-optimize for data-driven online inventory routing, *Transportation Science*, vol. 59, no. 5, pp. 1032-1056, 2025.
- Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q. and Liu, T., LightGBM: A highly efficient gradient boosting decision tree, *Advances in Neural Information Processing Systems*, vol. 30, pp. 3146-3154, 2017.

- Makridakis, S., Spiliotis, E. and Assimakopoulos, V., M5 accuracy competition: Results, findings, and conclusions, *International Journal of Forecasting*, vol. 38, no. 4, pp. 1346-1364, 2022.
- Mandi, J., Kotary, J., Berden, S., Mulamba, M., Bucarey, V., Guns, T. and Fioretto, F., Decision-focused learning: Foundations, state of the art, benchmark and future opportunities, *Journal of Artificial Intelligence Research*, vol. 80, pp. 1623-1701, 2024.
- Sethi, S. P. and Sorger, G., A theory of rolling horizon decision making, *Annals of Operations Research*, vol. 29, no. 1-4, pp. 387-416, 1991.
- Silver, E. A., Pyke, D. F. and Peterson, R., *Inventory Management and Production Planning and Scheduling*, 3rd Edition, John Wiley and Sons, New York, 1998.
- Stadtler, H., Kilger, C. and Meyr, H., *Supply Chain Management and Advanced Planning: Concepts, Models, Software, and Case Studies*, 5th Edition, Springer, 2015.
- Vlachos, I. and Reddy, P. G., Machine learning in supply chain management: Systematic literature review and future research agenda, *International Journal of Production Research*, vol. 63, no. 16, pp. 5987-6016, 2025.
- Walter, A., Ahsan, K. and Rahman, S., Application of artificial intelligence in demand planning for supply chains: A systematic literature review, *The International Journal of Logistics Management*, vol. 36, no. 3, pp. 672-719, 2025.